

# Technical Report: Age-Variant Face Matching System

## Abstract

This report presents a comprehensive age-variant face matching system capable of identifying the same individual across different ages. The system combines deep learning-based age prediction with sophisticated face matching techniques to handle the challenging task of recognizing faces with significant temporal variations. Our approach achieves robust performance across diverse age ranges and demographic groups through careful model design and training strategies.

## 1. Introduction

### 1.1 Problem Statement

Traditional face recognition systems struggle with age progression, as facial features change significantly over time due to aging processes. This research addresses the challenge of identifying the same person at different life stages by developing an integrated system that:

1. Predicts the age of individuals in facial images
2. Performs face matching with age-aware adjustments
3. Combines both predictions for robust age-variant identification

### 1.2 Objectives

- Develop accurate age estimation models for diverse demographics
- Implement robust face matching capable of handling temporal variations
- Create an integrated pipeline for age-variant face recognition
- Evaluate system performance across different age gaps and scenarios

## 2. Dataset Selection and Preprocessing

### 2.1 Primary Dataset: UTKFace

Rationale for Selection:

- **Comprehensive Age Range:** Covers ages 0-116 years, providing extensive temporal coverage
- **Demographic Diversity:** Includes multiple ethnicities and both genders
- **Large Scale:** Over 20,000 high-quality facial images
- **Consistent Format:** Standardized naming convention facilitating automated parsing

- **Real-world Variation:** Natural lighting, poses, and expressions

#### **Dataset Characteristics:**

- Total Images: 23,708
- Age Range: 0-116 years
- Gender Distribution: Male (11,357), Female (12,351)
- Ethnicity Distribution: White (10,078), Black (4,526), Asian (3,434), Indian (3,975), Others (1,695)
- Image Resolution: Variable (processed to 224×224)
- Format: [age]\_[gender]\_[race]\_[date&time].jpg

## **2.2 Backup Datasets**

#### **LFW (Labeled Faces in the Wild):**

- Used for face matching when UTKFace unavailable
- 13,000+ images of 5,749 people
- Real-world conditions with natural variations

## **2.3 Data Preprocessing Pipeline**

#### **Image Preprocessing:**

# Training Transform (Age Estimation)

```
train_transform = T.Compose([
    T.Resize((224, 224)),
    T.RandomHorizontalFlip(p=0.5),
    T.RandomRotation(degrees=15),
    T.ColorJitter(brightness=(0.5, 1.5), contrast=(0.5, 1.5)),
    T.RandomAffine(degrees=0, translate=(0.1, 0.1)),
    T.ToTensor(),
    T.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
```

# Evaluation Transform

```
eval_transform = T.Compose([
    T.Resize((224, 224)),
    T.ToTensor(),
    T.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
```

#### **Data Augmentation Strategy:**

- **Geometric:** Horizontal flips, rotations, affine transformations
- **Photometric:** Color jittering, brightness/contrast variations
- **Normalization:** ImageNet statistics for transfer learning compatibility

### Dataset Splitting:

- **Training:** 68% (stratified by gender/ethnicity)
- **Validation:** 12% (for hyperparameter tuning)
- **Testing:** 20% (for final evaluation)

## 3. Age Prediction Model Architecture

### 3.1 Model Design Philosophy

#### Architectural Choices:

- **Backbone Selection:** ResNet50
- **Transfer Learning:** Pre-trained ImageNet weights
- **Regression Head:** Custom multi-layer design
- **Regularization:** Dropout and batch normalization

### 3.2 ResNet50 Architecture

```
class AgeEstimationModel(nn.Module):
    def __init__(self, model_name='resnet', pretrain_weights='IMAGENET1K_V2'):
        # Backbone: ResNet50 with pre-trained weights
        self.backbone = resnet50(weights=ResNet50_Weights.IMAGENET1K_V2)

        # Custom regression head
        self.classifier = nn.Sequential(
            nn.Dropout(p=0.2),
            nn.Linear(2048, 512),
            nn.ReLU(inplace=True),
            nn.BatchNorm1d(512),
            nn.Dropout(p=0.2),
            nn.Linear(512, 256),
            nn.ReLU(inplace=True),
            nn.BatchNorm1d(256),
            nn.Dropout(p=0.2),
            nn.Linear(256, 1) # Single output for age
        )
```

#### Design Rationale:

- **Pre-trained Backbone:** Leverages learned features from ImageNet
- **Progressive Dimensionality Reduction:** 2048→512→256→1
- **Regularization:** Dropout (20%) and batch normalization
- **Activation:** ReLU for non-linearity

### 3.3 Loss Function Selection

### Primary Loss: L1 Loss (Mean Absolute Error)

```
loss_fn = nn.L1Loss()
```

#### Rationale:

- **Robustness:** Less sensitive to outliers than L2
- **Interpretability:** Direct age difference in years
- **Stability:** Consistent gradients across age ranges

#### Alternative Loss Functions Implemented:

- **Smooth L1 (Huber Loss):** Combines L1 and L2 benefits
- **Weighted MSE:** Emphasizes certain age ranges
- **Custom Age Loss:** Multi-component loss function

## 3.4 Training Strategy

#### Optimizer Configuration:

```
optimizer = optim.SGD(  
    model.parameters(),  
    lr=0.0001,  
    momentum=0.9,  
    weight_decay=0.001  
)
```

#### Learning Rate Scheduling:

- **StepLR:** Decay every 20 epochs by factor 0.1
- **CosineAnnealingLR:** Smooth decay with restarts
- **ReduceLROnPlateau:** Adaptive based on validation loss

#### Training Configuration:

- **Epochs:** 50 (with early stopping)
- **Batch Size:** 128
- **Early Stopping:** Patience of 10 epochs
- **Gradient Clipping:** Prevents exploding gradients

## 4. Face Matching Model Architecture

### 4.1 Siamese Network Design

#### Architecture Overview:

```
class SiameseNetwork(nn.Module):
```

```

def __init__(self, embedding_dim=512):
    self.face_embedding = FaceEmbeddingNetwork(embedding_dim)

def forward(self, img1, img2):
    emb1 = self.face_embedding(img1)
    emb2 = self.face_embedding(img2)
    return emb1, emb2

```

### Embedding Network:

```

class FaceEmbeddingNetwork(nn.Module):
    def __init__(self, embedding_dim=512):
        # ResNet50 backbone
        self.backbone = resnet50(pretrained=True)
        self.backbone.fc = nn.Identity()

        # Custom embedding layers
        self.embedding_layers = nn.Sequential(
            nn.Linear(2048, 1024),
            nn.BatchNorm1d(1024),
            nn.ReLU(inplace=True),
            nn.Dropout(0.5),
            nn.Linear(1024, embedding_dim),
            nn.BatchNorm1d(embedding_dim)
        )

    def forward(self, x):
        features = self.backbone(x)
        embeddings = self.embedding_layers(features)
        return F.normalize(embeddings, p=2, dim=1) # L2 normalization

```

## 4.2 Loss Function: Contrastive Loss

### Mathematical Formulation:

```

class ContrastiveLoss(nn.Module):
    def __init__(self, margin=1.0):
        self.margin = margin

    def forward(self, emb1, emb2, label):
        distance = F.pairwise_distance(emb1, emb2, p=2)

        loss_positive = label * torch.pow(distance, 2)
        loss_negative = (1 - label) * torch.pow(
            torch.clamp(self.margin - distance, min=0.0), 2
        )

```

```
return (loss_positive + loss_negative).mean()
```

#### Design Rationale:

- **Margin:** 1.0 provides good separation between classes
- **Euclidean Distance:** L2 distance in embedding space
- **Contrastive Nature:** Pulls similar pairs together, pushes dissimilar apart

### 4.3 Training Data Generation

#### Face Pair Creation:

```
def _create_pairs(self, pairs_per_identity):
    # Positive pairs (same person)
    for identity, images in identity_groups.items():
        if len(images) > 1:
            for i in range(len(images)):
                for j in range(i + 1, len(images)):
                    self.pairs.append((images[i], images[j]))
                    self.labels.append(1) # Same person

    # Negative pairs (different persons)
    for _ in range(num_positive):
        id1, id2 = random.sample(identities, 2)
        img1 = random.choice(identity_groups[id1])
        img2 = random.choice(identity_groups[id2])
        self.pairs.append((img1, img2))
        self.labels.append(0) # Different person
```

### 4.4 Distance Metrics and Thresholds

#### Primary Metrics:

- **Euclidean Distance:**  $\|emb1 - emb2\|_2$
- **Cosine Similarity:**  $(emb1 \cdot emb2) / (\|emb1\| \times \|emb2\|)$

#### Threshold Optimization:

- **Distance Threshold:** 0.5 (optimized on validation set)
- **Similarity Threshold:** 0.5 (standard cosine similarity cutoff)

## 5. Integrated Age-Variant Pipeline

### 5.1 Pipeline Architecture

### Main Components:

1. **Age Prediction Module:** Estimates ages for both images
2. **Face Matching Module:** Computes face similarity metrics
3. **Age-Variant Analysis:** Adjusts thresholds based on age difference
4. **Decision Fusion:** Combines evidence for final classification

## 5.2 Age-Aware Threshold Adjustment

### Aging Factor Calculation:

```
def _analyze_age_variant_matching(self, age_pred, face_match):
    age_diff = age_pred['age_difference']

    # Determine aging category and factor
    if age_diff <= 5:
        aging_category = 'minimal'
        aging_factor = 1.0
    elif age_diff <= 15:
        aging_category = 'moderate'
        aging_factor = 1.15
    elif age_diff <= 30:
        aging_category = 'significant'
        aging_factor = 1.3
    else:
        aging_category = 'extreme'
        aging_factor = 1.5

    # Adjust thresholds
    adjusted_distance_threshold = base_threshold * aging_factor
    adjusted_similarity_threshold = base_threshold / aging_factor
```

### Rationale:

- **Progressive Adjustment:** Larger age gaps require more lenient thresholds
- **Evidence-Based:** Categories based on facial aging research
- **Symmetric Adjustment:** Both distance and similarity metrics adjusted

## 5.3 Decision Fusion Strategy

### Evidence Collection:

```
evidence = {
    'basic_distance_match': face_match['same_person_distance'],
    'basic_similarity_match': face_match['same_person_similarity'],
    'age_aware_distance_match': age_analysis['age_aware_distance_match'],
    'age_aware_similarity_match': age_analysis['age_aware_similarity_match'],
    'age_compatibility_high': age_analysis['age_compatibility_score'] > 0.7,
```

```
'aging_probability_high': age_analysis['aging_probability'] > 0.5
}
```

#### Decision Logic:

- **Strong Evidence** ( $\geq 4/6$ ): Same person, high confidence
- **Moderate Evidence** ( $3/6$ ): Same person, medium confidence
- **Weak Evidence** ( $2/6$ ): Possibly same person, low confidence
- **Insufficient Evidence** ( $\leq 1/6$ ): Different person

## 5.4 Confidence Scoring

#### Multi-Factor Confidence:

```
def _calculate_confidence_scores(self, results):
    age_conf = mean(age_predictions_confidence)
    face_conf = mean(face_matching_confidence)
    age_compatibility = age_compatibility_score

    overall_confidence = min(age_conf, face_conf, age_compatibility)
    return overall_confidence
```

## 6. Performance Analysis and Evaluation

### 6.1 Age Prediction Performance

#### Evaluation Metrics:

- **MAE (Mean Absolute Error)**: Primary metric for age accuracy
- **RMSE (Root Mean Square Error)**: Sensitivity to large errors
- **Accuracy@k**: Percentage within k years ( $k=1,3,5,10$ )
- **R<sup>2</sup>**: Coefficient of determination

#### Expected Performance:

MAE: < 5 years

RMSE: < 7 years

Accuracy@5: > 75%

Accuracy@10: > 90%

R<sup>2</sup>: > 0.85

### 6.2 Face Matching Performance

#### Evaluation Metrics:



- **AUC (Area Under Curve):** Overall discriminative ability
- **Accuracy:** Classification accuracy at optimal threshold
- **Precision/Recall:** Class-specific performance
- **EER (Equal Error Rate):** Balanced error rate

**Expected Performance:**

AUC: > 0.85

Accuracy: > 80%

Precision: > 0.82

Recall: > 0.78

EER: < 0.15

## 6.3 Integrated System Performance

**Age-Variant Specific Metrics:**

- **Age-Aware Accuracy:** Performance considering age differences
- **Temporal Robustness:** Consistency across age gaps
- **Demographic Fairness:** Performance across ethnic groups

**Performance by Age Gap:**

0-5 years: > 90% accuracy

5-15 years: > 80% accuracy

15-30 years: > 70% accuracy

30+ years: > 60% accuracy

## 6.4 Evaluation Protocol

**Test Set Composition:**

- **Balanced Age Distribution:** Equal representation across age groups
- **Controlled Age Gaps:** Systematic evaluation across different temporal spans
- **Demographic Stratification:** Ensuring fair evaluation across groups

**Cross-Validation Strategy:**

- **Identity-Based Splits:** No identity overlap between train/test
- **Temporal Validation:** Separate evaluation for different age gaps
- **Demographic Validation:** Performance analysis per demographic group

# 7. System Capabilities and Limitations

## 7.1 Strengths

### Technical Strengths:

1. **Robust Age Prediction:** Handles ages 0-85+ with high accuracy
2. **Multi-Demographic Support:** Works across ethnicities and genders
3. **Age-Aware Matching:** Intelligent threshold adjustment
4. **Modular Design:** Flexible architecture supporting different backbones
5. **Comprehensive Evaluation:** Multiple metrics and confidence scoring
6. **Real-time Performance:** Optimized for practical deployment

### Methodological Strengths:

1. **Transfer Learning:** Leverages pre-trained models effectively
2. **Data Augmentation:** Robust training through synthetic variations
3. **Loss Function Diversity:** Multiple loss options for different scenarios
4. **Early Stopping:** Prevents overfitting through validation monitoring
5. **Confidence Estimation:** Provides reliability measures for predictions

## 7.2 Limitations and Challenges

### Technical Limitations:

1. **Extreme Age Gaps:** Performance degrades for differences >40 years
2. **Image Quality Sensitivity:** Requires reasonable resolution and lighting
3. **Pose Variation:** Limited robustness to extreme head poses
4. **Occlusion Handling:** Struggles with partially covered faces
5. **Dataset Bias:** Performance may vary with demographic representation

### Methodological Limitations:

1. **Threshold Sensitivity:** Performance dependent on threshold optimization
2. **Age Estimation Errors:** Cascade effects from incorrect age predictions
3. **Training Data Requirements:** Needs substantial paired data for face matching
4. **Computational Requirements:** High memory usage for large batch processing
5. **Generalization:** May not transfer well to different domains/populations

## 7.3 Failure Scenarios

### Common Failure Cases:

1. **Poor Image Quality:** Blurry, low-resolution, or poorly lit images
2. **Extreme Poses:** Profile views or unusual head orientations
3. **Significant Occlusions:** Masks, sunglasses, or other face coverings
4. **Unusual Expressions:** Extreme facial expressions altering facial geometry
5. **Cross-Domain Images:** Different image sources with varying characteristics
6. **Edge Demographics:** Underrepresented groups in training data
7. **Surgical Modifications:** Plastic surgery or other facial alterations
8. **Lighting Conditions:** Extreme shadows or overexposure

## 8. Implementation Details

### 8.1 Software Architecture

#### Modular Design:

```
age/                # Age estimation module
├── config/          # Configuration management
├── data/             # Data loading and preprocessing
├── models/           # Model architectures
├── training/         # Training utilities
├── inference/        # Inference pipeline
└── utils/           # Helper functions

face/               # Face matching module
├── config/          # Configuration management
├── data/             # Face pair creation
├── models/           # Siamese networks
├── training/         # Training pipeline
├── inference/        # Matching inference
└── utils/           # Visualization tools

main.py             # Integrated pipeline
```

### 8.2 Configuration Management

#### Age Model Configuration:

```
config = {
    'img_size': 224,
    'batch_size': 128,
    'lr': 0.0001,
    'epochs': 50,
    'model_name': 'resnet',
    'pretrain_weights': 'IMAGENET1K_V2',
    'loss_function': 'l1',
    'dropout_rate': 0.2,
    'early_stopping_patience': 10
}
```

#### Face Model Configuration:

```
{
    "embedding_dim": 512,
    "margin": 1.0,
    "learning_rate": 0.0001,
```

```
"batch_size": 32,  
"num_epochs": 30,  
"distance_threshold": 0.5,  
"pairs_per_identity": 5  
}
```

## 8.3 Memory Optimization

### Techniques Implemented:

```
# Gradient accumulation for large batch sizes  
accumulated_loss = 0  
for i, batch in enumerate(dataloader):  
    loss = model(batch) / accumulation_steps  
    loss.backward()  
    accumulated_loss += loss.item()  
  
    if (i + 1) % accumulation_steps == 0:  
        optimizer.step()  
        optimizer.zero_grad()  
  
# Memory cleanup  
del outputs  
torch.cuda.empty_cache()
```